

Navaneeth Maruthi

maruthi.n@northeastern.edu | (857)-421-9106 | [LinkedIn](#) | [Github](#) | [Portfolio](#)

EDUCATION

Northeastern University

Boston, MA

MS in Computer Science

Expected Graduation: May 2027

Relevant Coursework: Algorithms & Data Structures, Web Development, Building Scalable Distributed Systems, Database Management Systems, Machine Learning, Programming Design Paradigms

Visvesvaraya Technological University

Bangalore, India

BE in Computer Science and Engineering

May 2018 – July 2022

TECHNICAL SKILLS

Languages: Python, Java, JavaScript (ES6), SQL, C++, Go, C#

Frontend: React.js, Angular, HTML5, CSS3

Backend: Node.js, Express.js, Spring Boot, REST APIs, Apache Kafka, Apache Spark

Databases: MySQL, MongoDB (Atlas), PostgreSQL, Snowflake, Hive, NoSQL

Cloud & DevOps: AWS (S3, ECS, ECR, Bedrock), Docker, Kubernetes, Terraform, GCP, CI/CD pipelines

Tools: Git, GitHub, Swagger, JWT, Agile/Scrum, Maven

Certifications: JPMorgan Chase Software Engineering Simulation — Forage (Feb 2026)

WORK EXPERIENCE

Cognizant

Nov 2022 – Oct 2024

Junior Software Engineer

Bangalore, India

- Architected a production-grade **Apache Spark** pipeline migrating **10+ years** of historical enterprise data across a **4-zone processing architecture** (Confirmed → Semantic → Provisioning → Refined), enforcing **data quality** and semantic consistency at each zone transition.
- Designed a **REST API-driven** pipeline registration system via a **Swagger-enabled interface**, enabling dynamic dataset onboarding and reusable pipeline configuration — eliminating hardcoded logic across multiple data domains.
- Refactored a **monolithic data ingestion workflow** into modular, reusable **Apache Spark** components, decoupling an always-on batch processing layer from on-demand API execution and reducing end-to-end pipeline overhead by **~40%**.
- Migrated **3+ production pipeline configurations** during an infrastructure consolidation, handling dataset re-registration, **IAM role** updates, and orchestration reconfiguration under **Control-M** with zero data loss.
- Delivered **Snowflake** as the analytical terminal sink of the migration pipeline, enforcing **schema contracts** and data quality standards that ensured reliable consumption by enterprise reporting and downstream analytics teams.
- Built a fault-tolerant **AWS S3** intermediate storage layer validated end-to-end using the **Prana** testing framework and **AWS Athena** across Refined and Trusted zones — ensuring structural integrity and semantic accuracy pre- and post-migration.

Cognizant

Apr 2022 – Sept 2022

Software Engineer Intern

Bangalore, India

- Built and shipped **SPA features** for an enterprise internal web application serving **100+ users** using **Spring Boot (Java)** and **Angular**, delivering UI components and backend endpoints across **Agile** sprint cycles with unit test coverage via Java testing frameworks.
- Integrated **Spring Boot REST APIs** into an existing platform to automate internal issue resolution workflows, reducing monthly support tickets by **30%** and removing manual intervention from recurring operational tasks.

PROJECTS

TradeFlow — Distributed FIX Protocol Engine ([Link](#))

Spring 2026

- Engineered the **Order Router** and **cost-aware circuit breaker** in **Go** for a distributed **FIX protocol** trading engine, routing orders across **5 AWS SQS queues** by financial impact and tripping when slippage exceeded **\$500/min** (CS 6650: Scalable Distributed Systems).
- Deployed **4 containerized microservices** to **AWS ECS Fargate** via **Terraform** IaC; instrumented **CloudWatch** for latency and cost tracking, and validated throughput under load with **Locust**.

Wisconsin Breast Cancer Diagnosis ([Link](#))

Dec 2025

- Built and compared 4 ML models (Logistic Regression, Random Forest, Decision Tree, K-NN) on the Wisconsin Diagnostic dataset; achieved **98.25% accuracy** and **0.997 AUC** via GridSearch tuning with a **4.76% false negative rate** — feature importance aligned with clinical oncological markers (radius_worst, concave_points_worst), confirming medically meaningful predictions.

Alignify ([Link](#))

Nov 2025

- Built a full-stack goal-planning app using **React**, **Node.js**, **Express**, and **MongoDB Atlas**, enabling **50+ users** to convert goals into structured action plans with **JWT authentication** and persistent end-to-end workflows.
- Containerized with **Docker** and deployed via **Render**; resolved **Lighthouse** contrast failures to achieve **WCAG-compliant** design across all components.